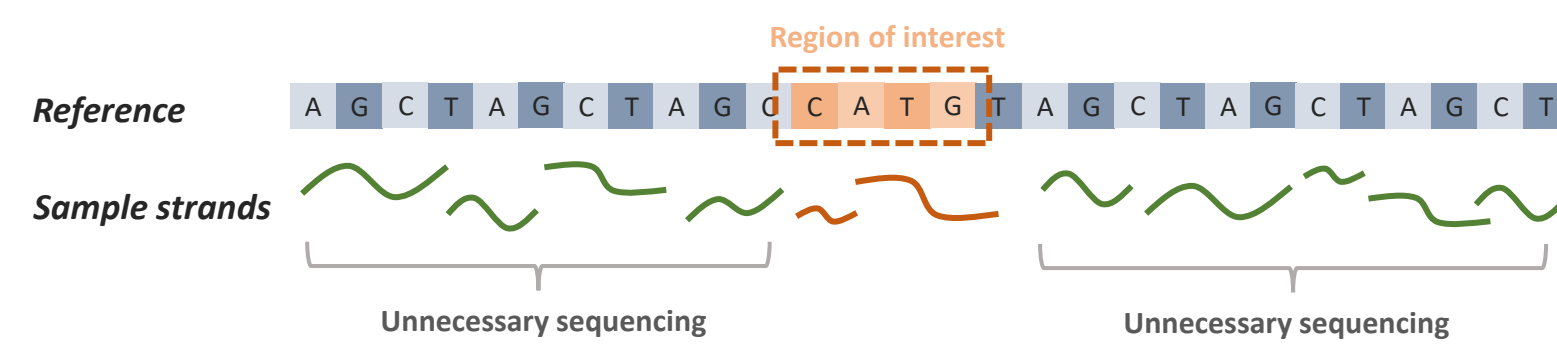


## Introduction

### Selective sequencing with nanopore technology enables efficient targeted genome analysis

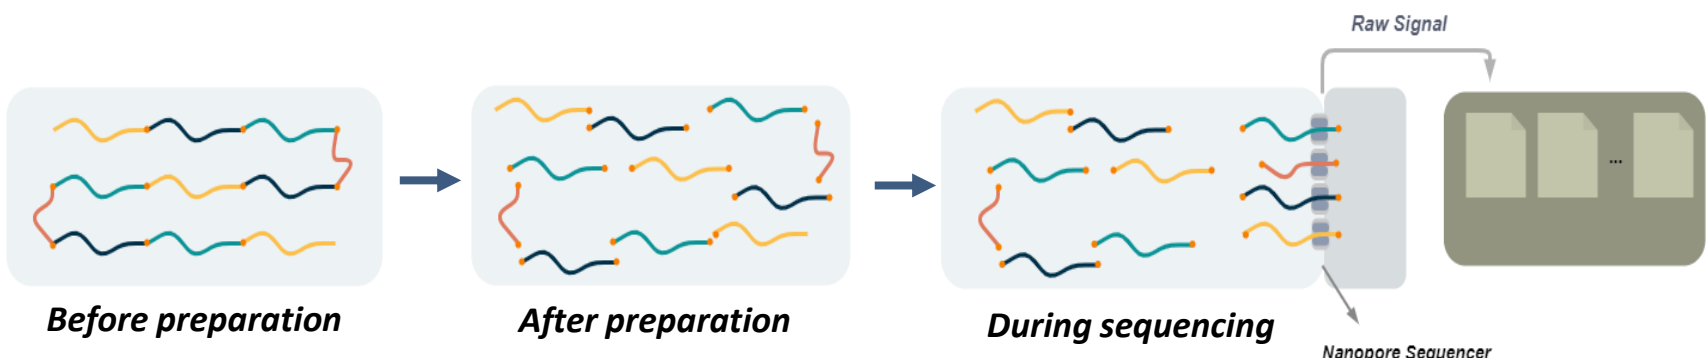


→ Fully **sequence** strands within **regions of interest**; **reject** others

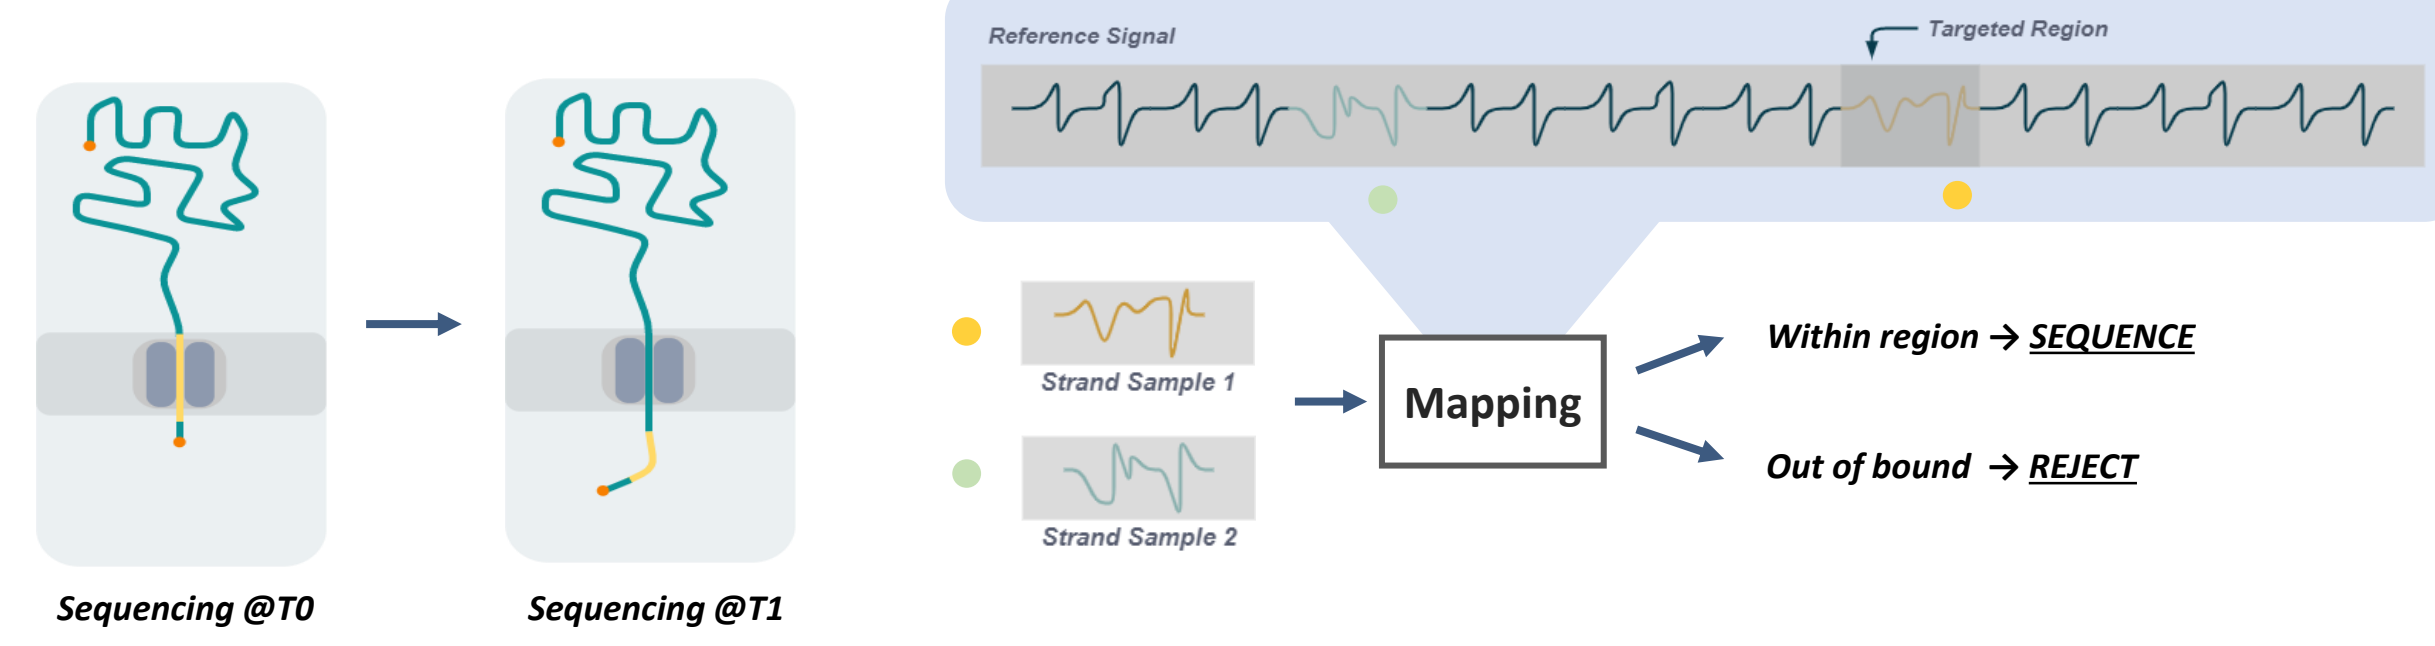
“Needle in a haystack”  
e.g. [regions of interest / others]  
→ [pathogen / host]  
→ [cancer / nontumor]

### Nanopore Sequencing

- Simple sample preparation (minimal priori knowledge)
- Real-time data output → real-time analysis
- Able to reject strands at individual nanopore channels

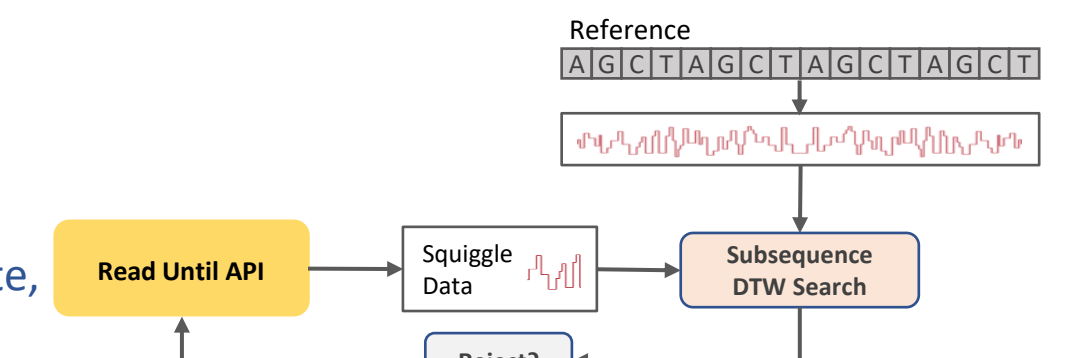


### Read Until

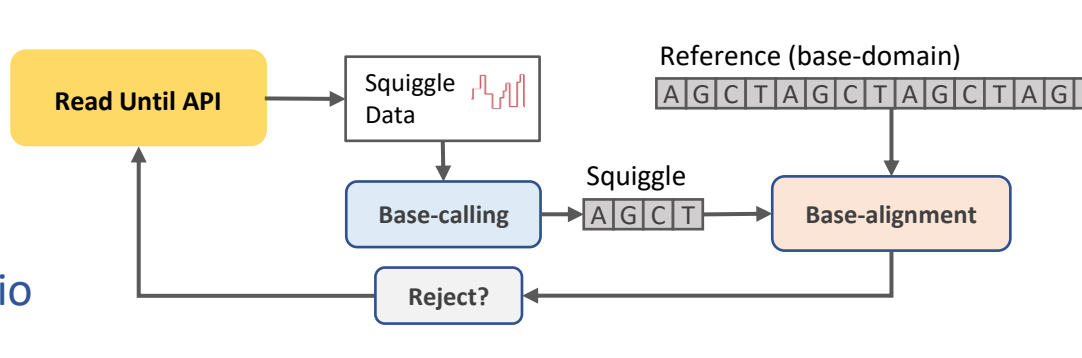


### Read Until Implementations

- Squiggle-domain sequence matching
  - First Read Until implementation (RUScripts [1])
  - Uses **Dynamic Time Warping (DTW)** to map sequences
  - Needed 22-core server to keep up with slower sequencing rate, deprecated after sequencing rate ↑



- Base-domain sequence matching
  - Base-calls** the squiggle and uses **base-alignment** to map [2]
  - Able to scale to giga-base references
  - Requires high-end GPU to perform real-time base-calling (excessive), loses portability, and has high performance-watt ratio



### HARU: The proposed Read Until implementation

- First FPGA accelerated Read Until implementation
- Software-hardware co-design targeting **low-cost** MPSoCs
- Extends the MinION sequencer’s **portable** nature
- Low performance requirement for host machine

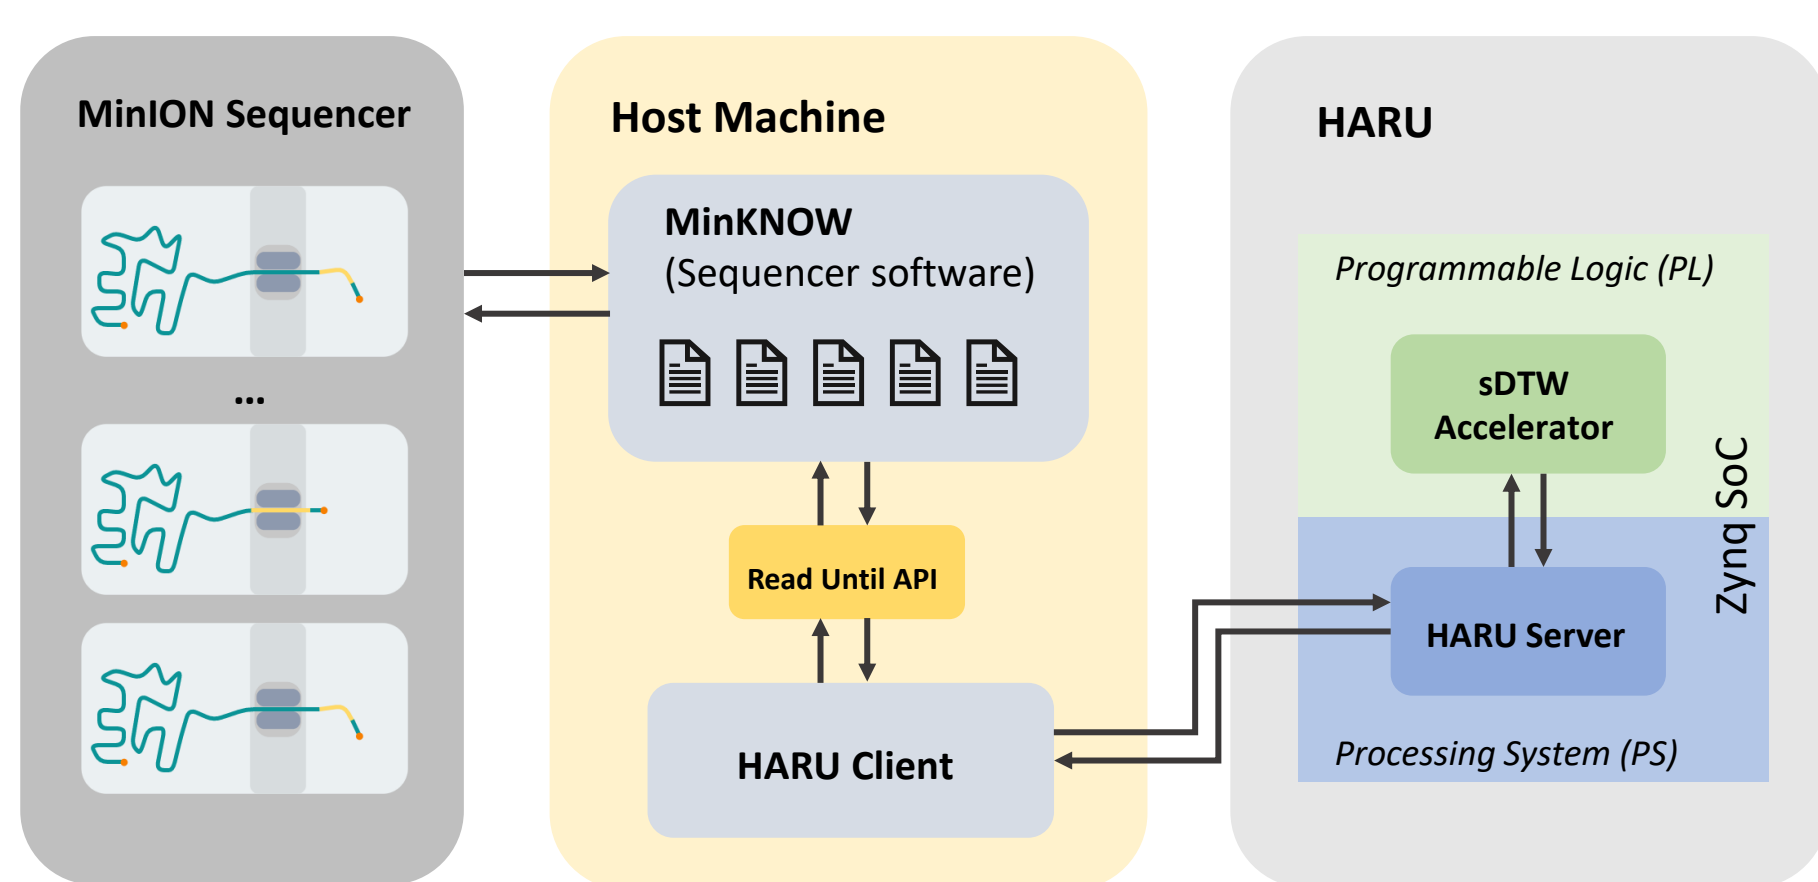


### Contributions

- Minimal requirements for performing targeted small-genome analysis
- Demonstrates the use of High-Level-Synthesis (HLS) for DNA sequencing and analysis acceleration
- Provides an extendible framework for Read Until

## HARU: Hardware Accelerated Read Until

### Read Until with HARU (MinION + HARU)



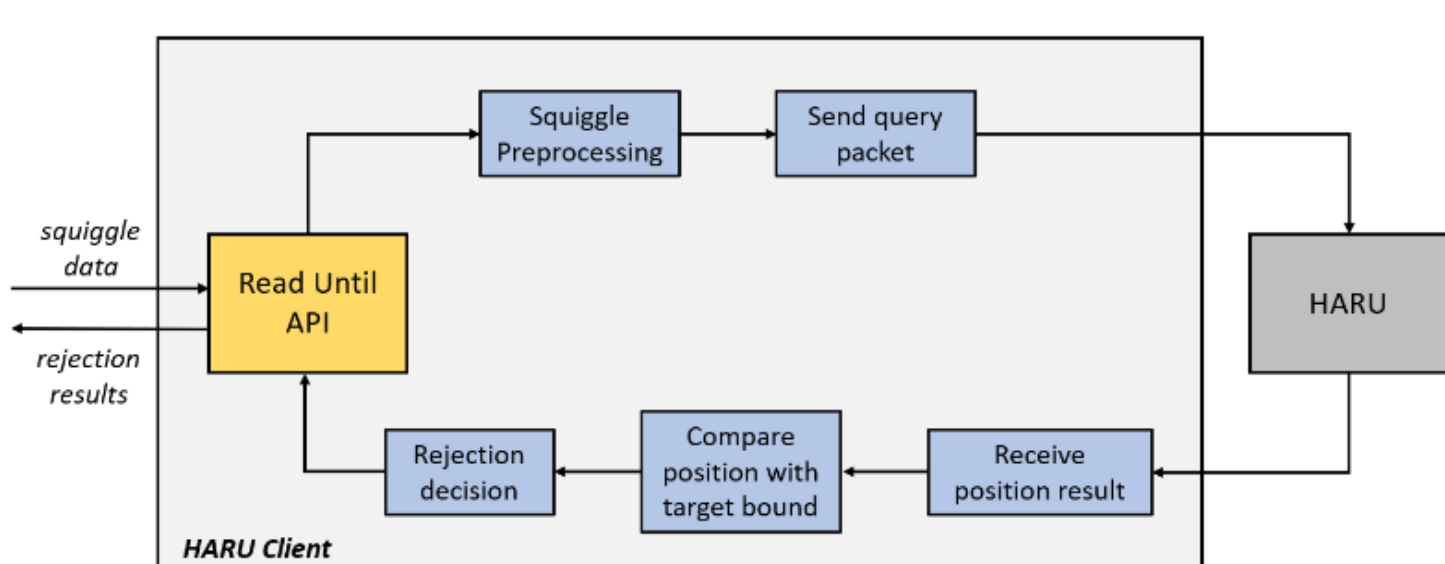
### HARU Client

#### Sequencer → HARU

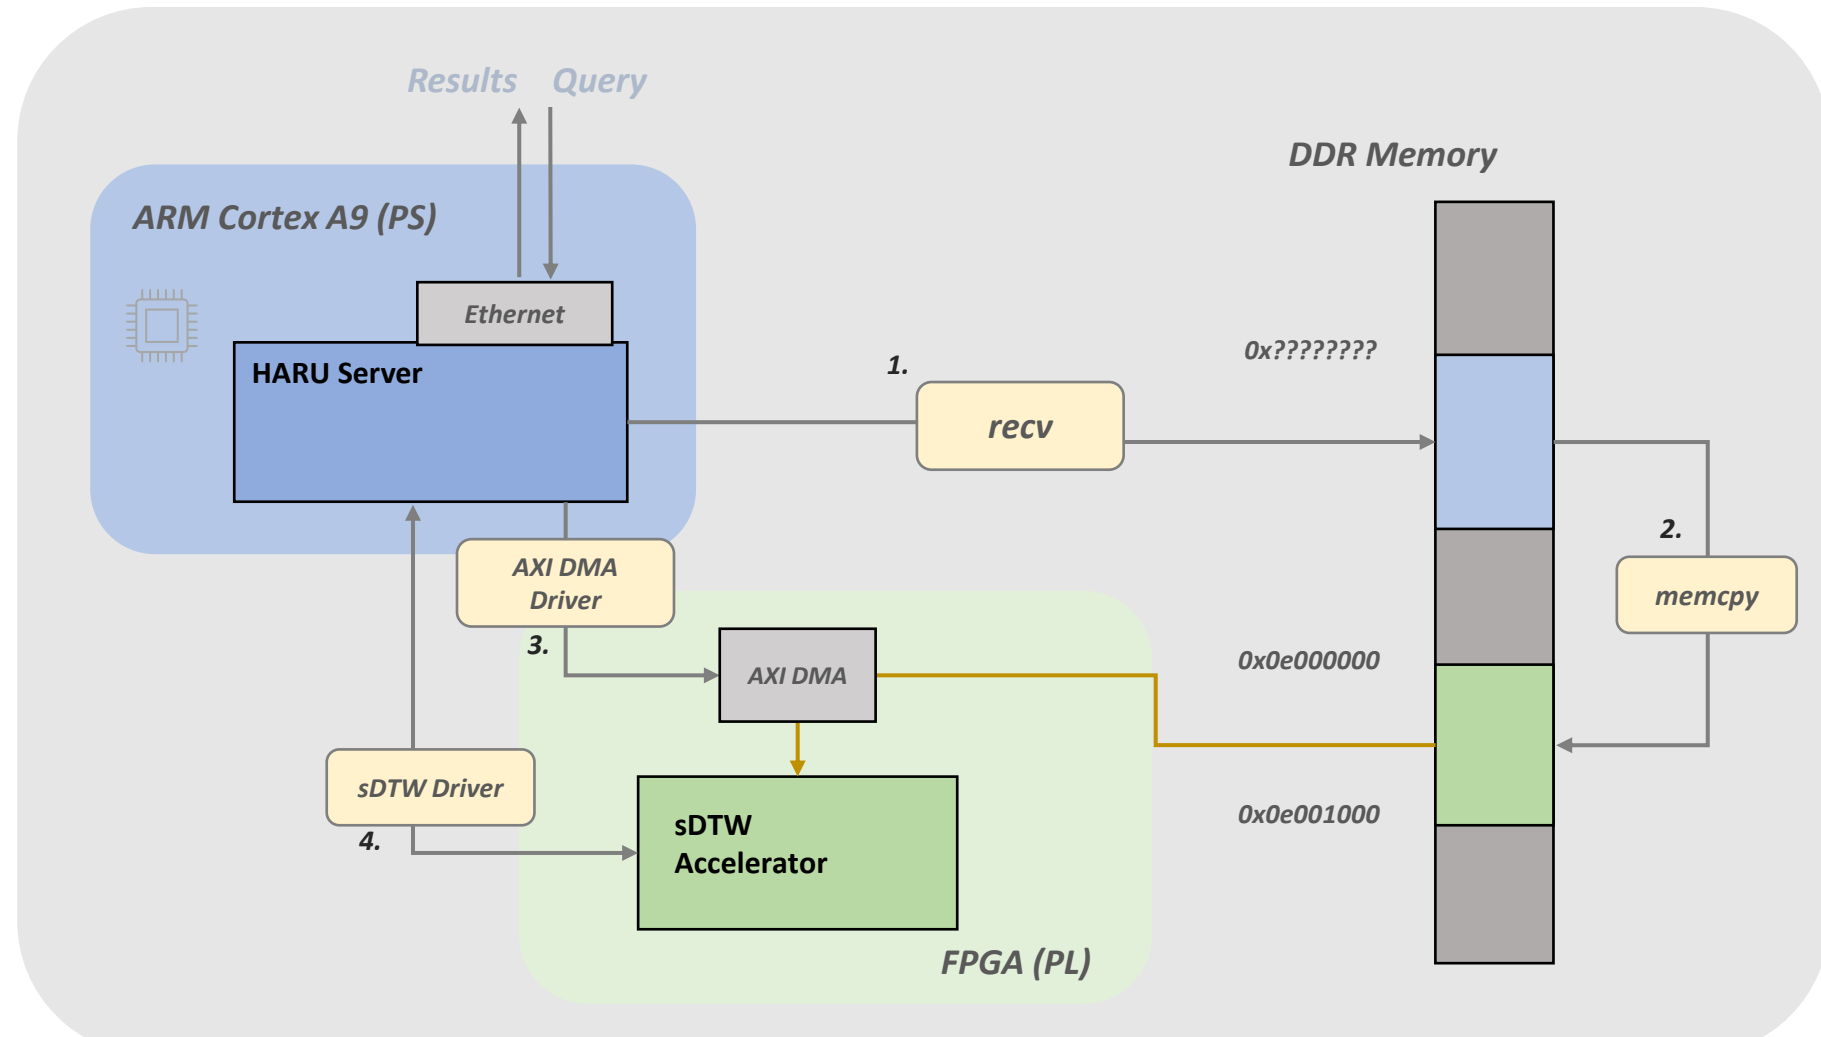
- Collects real-time squiggle data via the Read Until API
- Pre-processes raw data
- Sends data to HARU via Ethernet

#### HARU → Sequencer

- Receives sequence mapping results from HARU via Ethernet
- Determines whether the position of strand is within a region of interest
- Sends back rejection to sequencer software if not a necessary strand



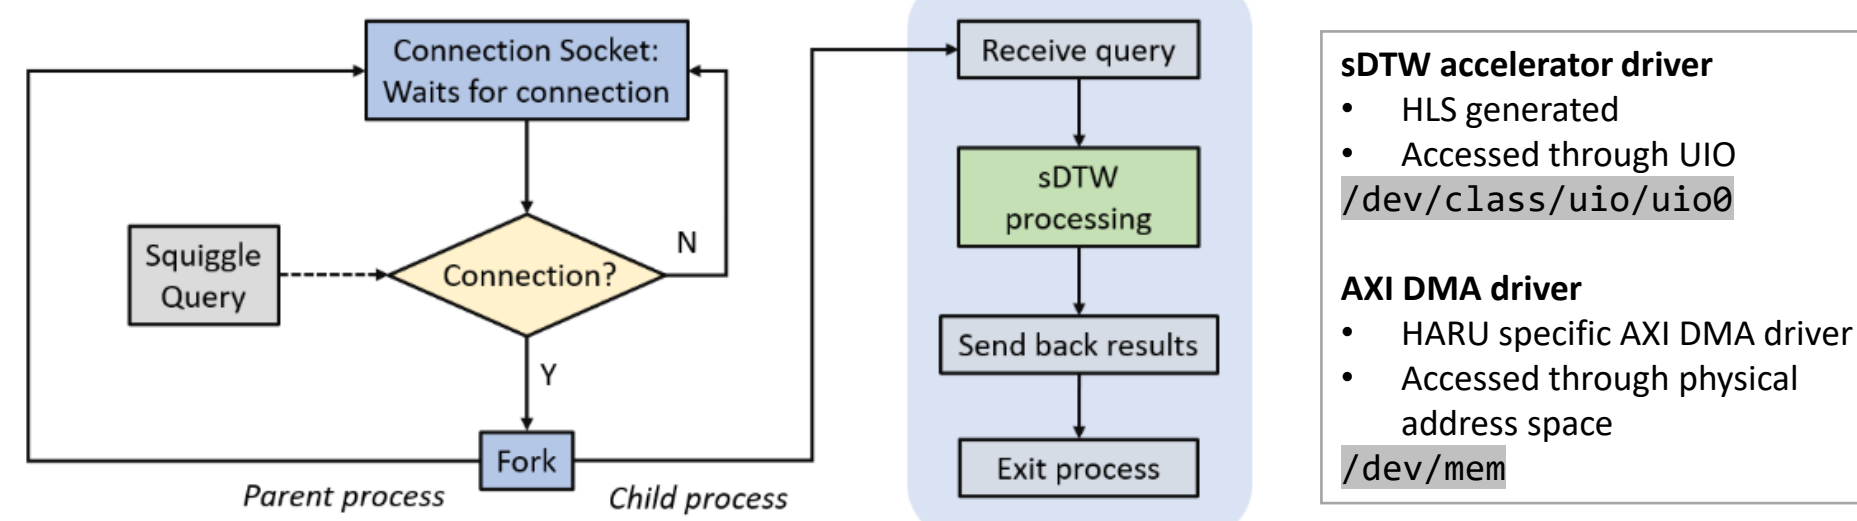
### HARU Overview



### HARU Server

- Server application running on a custom PetaLinux generated embedded Linux OS on the processing system of the Zynq MPSoC
- Responsible for query request handling
- Sends query over to accelerator via AXI stream (HP AXI)
- Controls the custom sDTW accelerator through custom drivers
- Sends results back to client via Ethernet using the same socket

**Documented AXIs throughput:**  
• MM2S = 399.04 MB/s  
• S2MM = 298.59 MB/s  
**Benchmarked AXIs throughput:**  
• 333.16 MB/s



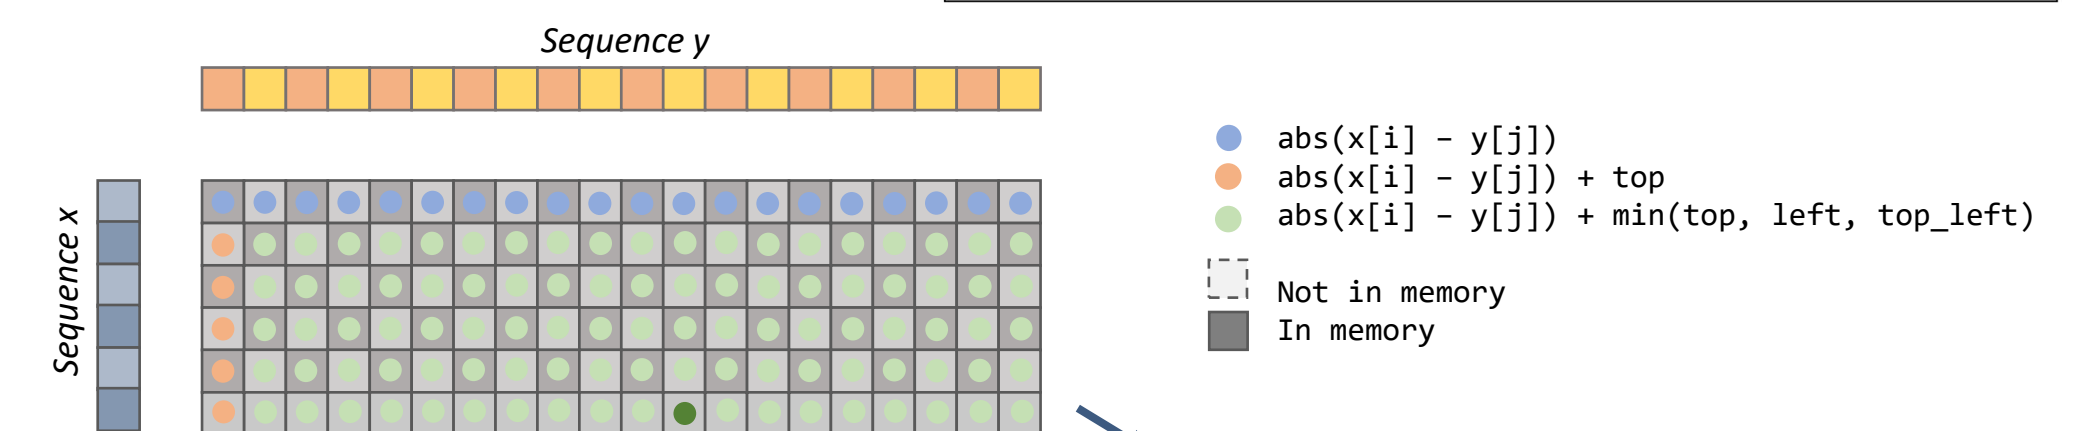
### Subsequence DTW Accelerator

#### Original subsequence DTW algorithm [3]:

- Given two sequences  $X, Y$
- $X := (x_1, x_2, \dots, x_M)$  of length  $M \in \mathbb{N}$
  - $Y := (y_1, y_2, \dots, y_N)$  of length  $N \in \mathbb{N}$

- and cost matrix  $C \in \mathbb{R}^{M \times N}$
- $C(m, n) := |x_m - y_n|$

**Algorithm: SUBSEQUENCE DTW** Exercise 7.6 from [Müller, FMP, Springer 2018]  
**Input:** Cost matrix  $C$  of size  $N \times M$   
**Output:** Accumulated cost matrix  $D$   
Indices  $a^*, b^* \in [1:M]$  of an optimal subsequence of  $Y$   
Optimal warping path  $P^*$  between  $X$  and  $Y(a^*: b^*)$   
**Procedure:** Initialize  $(N \times M)$  matrix  $D$  by  $D(n, 1) = \sum_{k=1}^n C(k, 1)$  for  $n \in [1:N]$  and  $D(1, m) = C(1, m)$  for  $m \in [1:M]$ . Then compute in a nested loop for  $n = 2, \dots, N$  and  $m = 2, \dots, M$ :  
 $D(n, m) = C(n, m) + \min(D(n-1, m-1), D(n-1, m), D(n, m-1))$ .  
Set  $b^* = \text{argmin}_{n \in [1:N]} D(n, M)$ . (If ‘argmin’ is not unique, take smallest index.)  
Set  $l = 1$  and  $q_l = (N, b^*)$ .  
Then repeat the following steps until  $q_l = (1, m)$  for some  $m \in [1:M]$ :  
Increase  $l$  by one and let  $(n, m) = q_{l-1}$ .  
If  $m = 1$ , then  $q_l = (n-1, 1)$ .  
else  $q_l = \text{argmin}_{n \in [1:N]} \{D(n-1, m-1), D(n-1, m), D(n, m-1)\}$ .  
(If ‘argmin’ is not unique, take lexicographically smallest cell.)  
Set  $l = l$  and  $a^* = m$ . Return  $D, a^*, b^*$ , and  $P^* = (q_l, q_{l-1}, \dots, q_1)$ .



#### Algorithmic optimisations:

- Padding for cost matrix
- Reduce cost matrix to single column

#### HLS specific optimisations:

- Pipelining of column computation
- 16-bit fixed point data type

#### Resulting accelerator:

- Oblique PE array with size of  $M$
- Cost matrix with size of  $3M$
- Oblique PE array propagates through the reference sequence

→ **Task latency** :=  $(M + N - 1) \times \text{Initiation Interval}$

## Results and Evaluation

### Experiment Details and Results

- Accelerator synthesised using Vivado HLS
- Targets the Xilinx Zynq-7020 device (xc7z020clg484-1)
- Tested on the target enrichment application for the bacteriophage lambda DNA
- Single direction has 48,502 bp, giving a full search space of 97,004 bp

#### Synthesis Results

	Slice LUTs	Slice Register	Slice	BRAM
Available (Zynq-7020)	53,200	106,400	13,300	140
HARU	32,341 (60.79%)	18,899 (17.76%)	9,615 (72.29%)	32.5 (23.21%)

#### HLS Latency Estimates

	Cycles	Clock Freq.	Estimated Time
Single directional reference search	48875	90 MHz	0.543 ms
Bi-directional reference search (Zynq-7020)	97755	90 MHz	1.086 ms
Unpack Streamed Query	250	90 MHz	2.778us
Overall Subseek DTW	98005	90 MHz	1.089 ms

#### Comparison with RUScripts

	RUScripts (reference)		HARU (proposed)		
	Laptop Intel i7-8565U	Desktop Intel i9-10850K	HARU system	Network latency	Overall latency
Avg. sDTW task latency	345.75 ms	136.11 ms	1 ms	3.36 ms	4.36 ms

#### Key results:

- Core sDTW: 345.75x faster than Intel i7 Laptop, 136.11x faster than Intel i9 Desktop
- Overall: 79.3x faster than RUScripts on Intel i7 Laptop, 31.22x faster than RUScripts on Intel i9 Desktop
- Bottleneck is now the network latency (currently unoptimized)

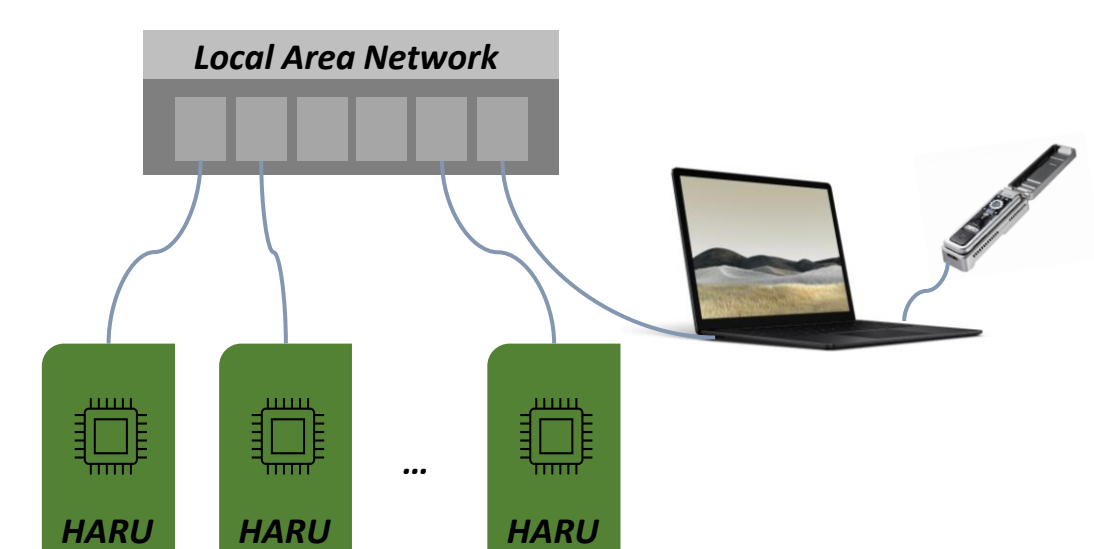
### Evaluation

#### Substantial speedup at a low hardware cost

- Subsequence DTW search now linearly dependant to the length of reference sequence
- Cost matrix only requires three times the size of squiggle sequence (subsequence)
- Optimal for smaller genomes (e.g. bacteria, virus)  
→ fast and direct search, can fully store the reference in on-chip memory (no sw-hw transfer overhead)

#### Preserves portability while enabling scalability

- Accesses HARU’s service through Ethernet
- No harsh requirements for host machine running HARU client
- Scalable by deploying a cluster of MPSoCs running HARU  
→ In-the-field analysis with low hardware requirements



#### Provides an extendible low-cost yet high performance-per-watt framework

- HARU demonstrated the use of HLS tools to perform acceleration for DNA sequencing and analysis techniques
- The framework is interchangeable and extendable based on application and algorithmic requirements

## References

- [1] M. Loose, S. Malla, and M. Stout, “Real time selective sequencing using nanopore technology,” BioRxiv, 2016.
- [2] A. Payne, N. Holmes, T. Clarke, R. Munro, B. Debebe, and M. W. Loose, “Nanopore adaptive sequencing for mixed samples, whole exome capture and targeted panels.,” BioRxiv, 2020.
- [3] M. Müller, “Dynamic time warping,” Information retrieval for music and motion, pp. 69–84, 2007.